

# Camila Blank

✉ camilab@stanford.edu    ⓧ camilablank.com    ⚙ camila-blank    ♻ camilablank

## Education

**Stanford University, BS in Mathematics and Computer Science**, GPA: 3.9/4.0.

- Courses: NLP w/ Deep Learning, Algorithms, Real Analysis, Linear Algebra, Abstract Algebra, Probability, Explainable AI, Ethical Theory, etc.

## Experience

### Brain Foundation Model Researcher

*Stanford Artificial Intelligence Laboratory*

*Stanford, CA*

*September 2025-Present*

- Building an explainable, end-to-end foundation model for diffusion MRI
- Analyzed UMAP embeddings to discover that model's latent space encodes clinically meaningful structure
- Mentored by Ehsan Adeli in STAI/SVL Lab, co-authored paper under review for CVPR 2026.

### ML and Computational Neuroscience Researcher

*The Linderman Lab at Stanford University*

*Stanford, CA*

*Jun 2025 - Present*

- Led project developing reinforcement learning algorithms to model intrinsic rewards in complex animal decision-making
- Awarded 2025 SURP-Stats fellowship, mentored by Scott Linderman and Aditi Jha

### Problem Writer

*Stanford Math Tournament*

*Stanford, CA*

*Sep 2024 – Present*

- Writing combinatorics problems for one of the largest university-run math contests globally

### Neuroscience Researcher

*The Gibson Lab at Stanford University School of Medicine*

*Stanford, CA*

*Jan 2025 – Jun 2025*

- Applied machine learning algorithms to analyze EEG sleep recording data
- Project investigated the role of myelination and circadian rhythms in Alzheimer's Disease

### Composer Fellow

*Los Angeles Philharmonic*

*Los Angeles, CA*

*Oct 2023 - Jun 2024*

- Composed 3 original works performed and recorded by the LA Phil, LA Master Chorale, and other ensembles

## Machine Learning Projects

### Model Organisms Of Unfaithful Chain-of-Thought

[Blog ↗](#)

- Developed method using Synthetic Document-Fine-Tuning to create a model organism of unfaithful CoT
- Discovered linear direction that can be steered/ablated to enhance or suppress unfaithfulness

### Exploring Dynamic Profiles of User Emotions in LLMs

[Blog ↗](#)

- Extracted hidden states from Llama-3.1-8B-Instruct and trained linear logistic probes
- Found emotion classification accuracy was >3.5x chance, updated promptly across conversations
- Steered responses using control-probe vectors, confirmed with LLM judge

### Reinforcement Learning Models of Animal Curiosity

[Poster ↗](#)

- Modeled mice learning to navigate 127-node binary tree maze
- Built Q-learning variants to model reward parameterizations and evolving world models in mice
- Custom algorithm using KL-divergence-based reward outperformed vanilla Q-learning in log-likelihood

## Skills

**Languages:** Spanish (Native proficiency), Mandarin Chinese (Professional working proficiency)

**Technical:** Python (Pytorch, Gymnasium, NumPy, Pandas, Matplotlib), C++, R, LaTeX